

# Chapter 7

## Multimedia Networking

---


### *A note on the use of these ppt slides:*

*We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a lot of work on our part. In return for use, we only ask the following:*

- ❖ *If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)*
- ❖ *If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.*

*Thanks and enjoy! JFK/KWR*

© All material copyright 1996-2016  
J.F Kurose and K.W. Ross, All Rights Reserved



## Computer Networking: A Top Down Approach

7<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson/Addison Wesley  
April 2016

# *Multimedia networking: outline*

*7.1 multimedia networking applications*

*7.2 streaming stored video*

*7.3 voice-over-IP*

# Multimedia networking: outline

*7.1 multimedia networking applications*

*7.2 streaming stored video*

*7.3 voice-over-IP*


# Multimedia: bit rate comparison

- ❖ *Frank is looking at Facebook pages*
  - *new photo every 10 seconds,*
  - *photos are on average 200 Kbytes in size*
- ❖ *Martha is listening to many MP3 songs*
  - *one after the other, each encoded at a rate of 128 kbps*
- ❖ *Victor is watching a video*

	Bit rate	Bytes transferred in 67 min
<b>Facebook Frank</b>	160 kbps	80 Mbytes
<b>Martha Music</b>	128 kbps	64 Mbytes
<b>Victor Video</b>	2 Mbps	1 Gbyte

# Multimedia: audio

- ❖ analog audio signal sampled at constant rate
  - telephone: 8,000 samples/sec
  - CD music: 44,100 samples/sec
- ❖ each sample quantized, i.e., rounded
  - e.g.,  $2^8=256$  possible quantized values
  - each quantized value represented by bits, e.g., 8 bits for 256 values




# Multimedia: audio

- ❖ example: 8,000 samples/sec, 256 quantized values: 64,000 bps
- ❖ receiver converts bits back to analog signal:
  - some quality reduction

## example rates

- ❖ CD: 1.411 Mbps
- ❖ MP3: 96, 128, 160 kbps



# Multimedia: video


- ❖ video: sequence of images displayed at constant rate
  - e.g. 24, 30 images/sec
- ❖ digital image: array of pixels
  - each pixel represented by bits
- ❖ coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (purple) and number of repeated values ( $N$ )



*frame  $i$*

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



*frame  $i+1$*

# Multimedia: video


- ❖ **CBR: (constant bit rate):**  
video encoding rate fixed
- ❖ **VBR: (variable bit rate):**  
video encoding rate changes as amount of spatial, temporal coding changes
- ❖ **examples:**
  - MPEG 1 (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

**spatial coding example:** instead of sending  $N$  values of same color (all purple), send only two values: color value (purple) and number of repeated values ( $N$ )



frame  $i$

**temporal coding example:**  
instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$



# Youtube bit rate

## Velocità in bit video consigliate per i caricamenti SDR

Tipo	Velocità in bit video, frequenza fotogrammi standard (24, 25, 30)	Velocità in bit video, frequenza fotogrammi elevata (48, 50, 60)
2160p (4K)	35-45 Mbps	53-68 Mbps
1440p (2K)	16 Mbps	24 Mbps
1080p	8 Mbps	12 Mbps
720p	5 Mbps	7,5 Mbps
480p	2,5 Mbps	4 Mbps
360p	1 Mbps	1,5 Mbps

*SDR: Standard Dynamic Range*  
*HDR: High Dynamic Range*

## Velocità in bit video consigliate per i caricamenti HDR

Tipo	Velocità in bit video, frequenza fotogrammi standard (24, 25, 30)	Velocità in bit video, frequenza fotogrammi elevata (48, 50, 60)
2160p (4K)	44-56 Mbps	66-85 Mbps
1440p (2K)	20 Mbps	30 Mbps
1080p	10 Mbps	15 Mbps
720p	6,5 Mbps	9,5 Mbps
480p	Non supportato	Non supportato
360p	Non supportato	Non supportato

## Velocità in bit audio consigliate per i caricamenti

Tipo	Velocità in bit audio
Mono	128 kbps
Stereo	384 kbps
5.1	512 kbps

**Source:** <https://support.google.com/youtube/answer/1722171?hl=it#zippy=%2Ccontentore-mp%2Ccodec-audio-aac-lc%2Cfrequenza-fotogrammi%2Cvelocit%C3%A0-in-bit>

# Current situation of mobile traffic



# Multimedia networking: 3 application types

- ❖ *streaming, stored* audio, video
  - *streaming*: can begin play out before downloading entire file
  - *stored (at server)*: can transmit faster than audio/video will be rendered (implies storing/buffering at client)
  - e.g., YouTube, Netflix
- ❖ *conversational* voice/video over IP
  - interactive nature of human-to-human conversation limits delay tolerance
    - $<150$ ,  $[150:400]$ ,  $>400$  ms
  - e.g., Skype
- ❖ *streaming live* audio, video
  - e.g., live sporting event (football)

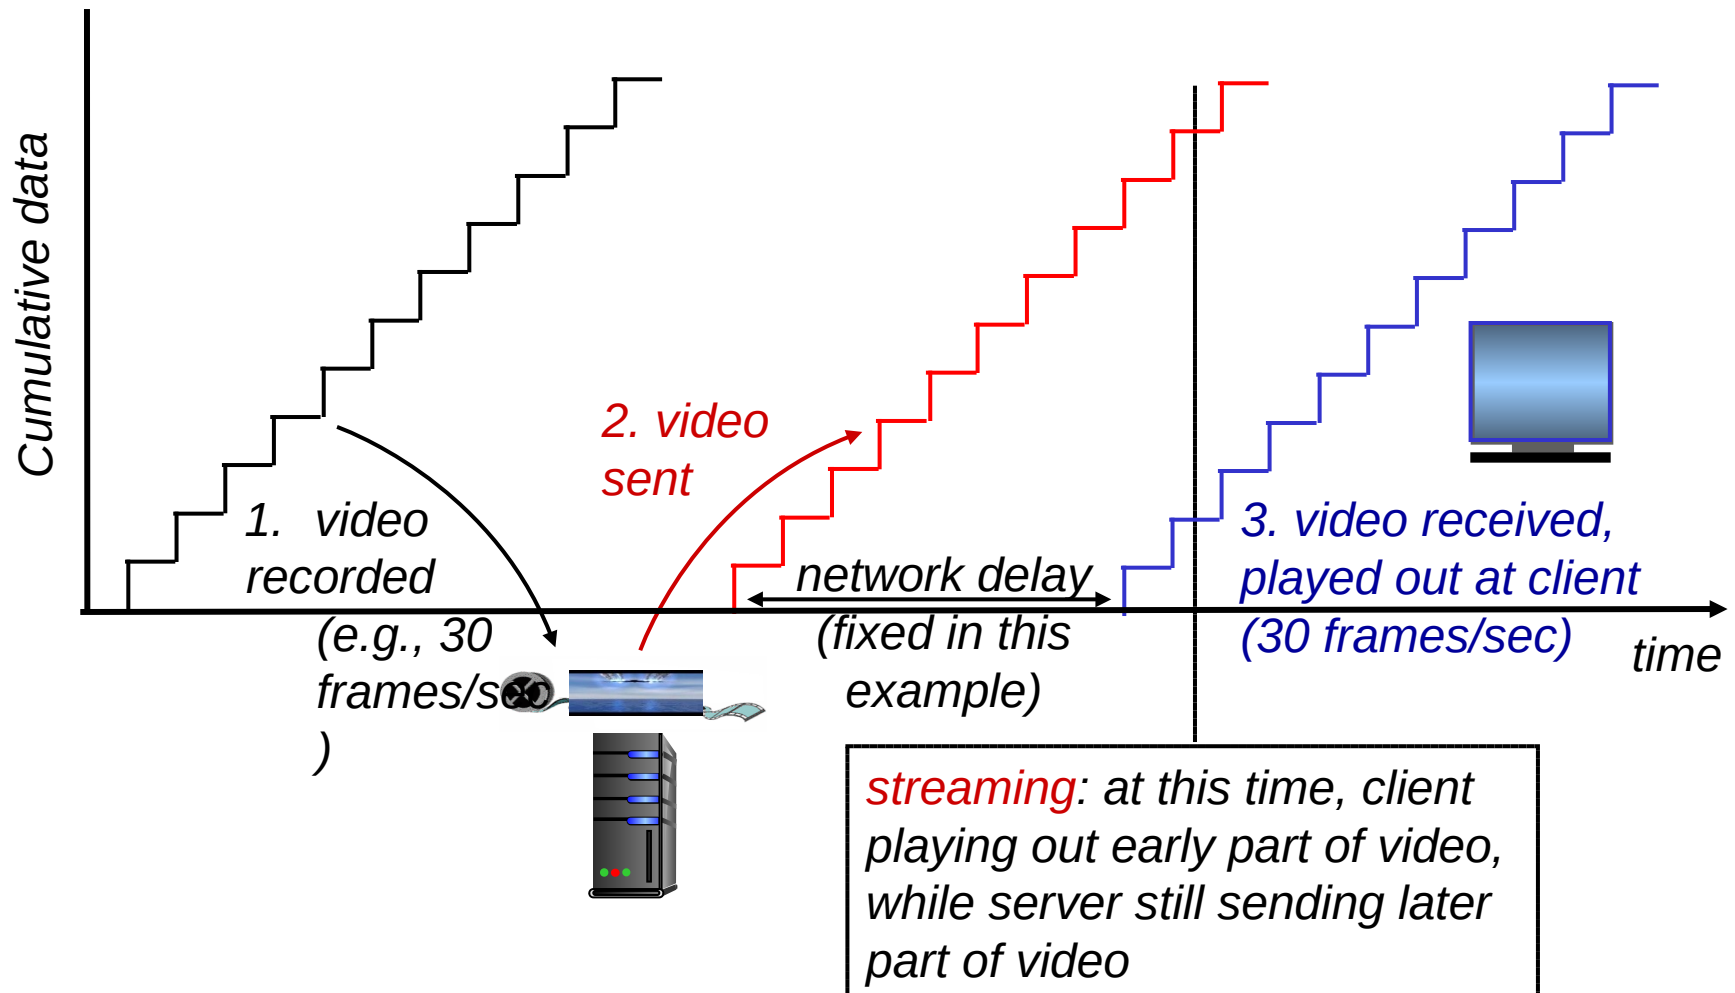
# *Multimedia networking: outline*

*7.1 multimedia networking applications*

*7.2 streaming stored video*

*7.3 voice-over-IP*

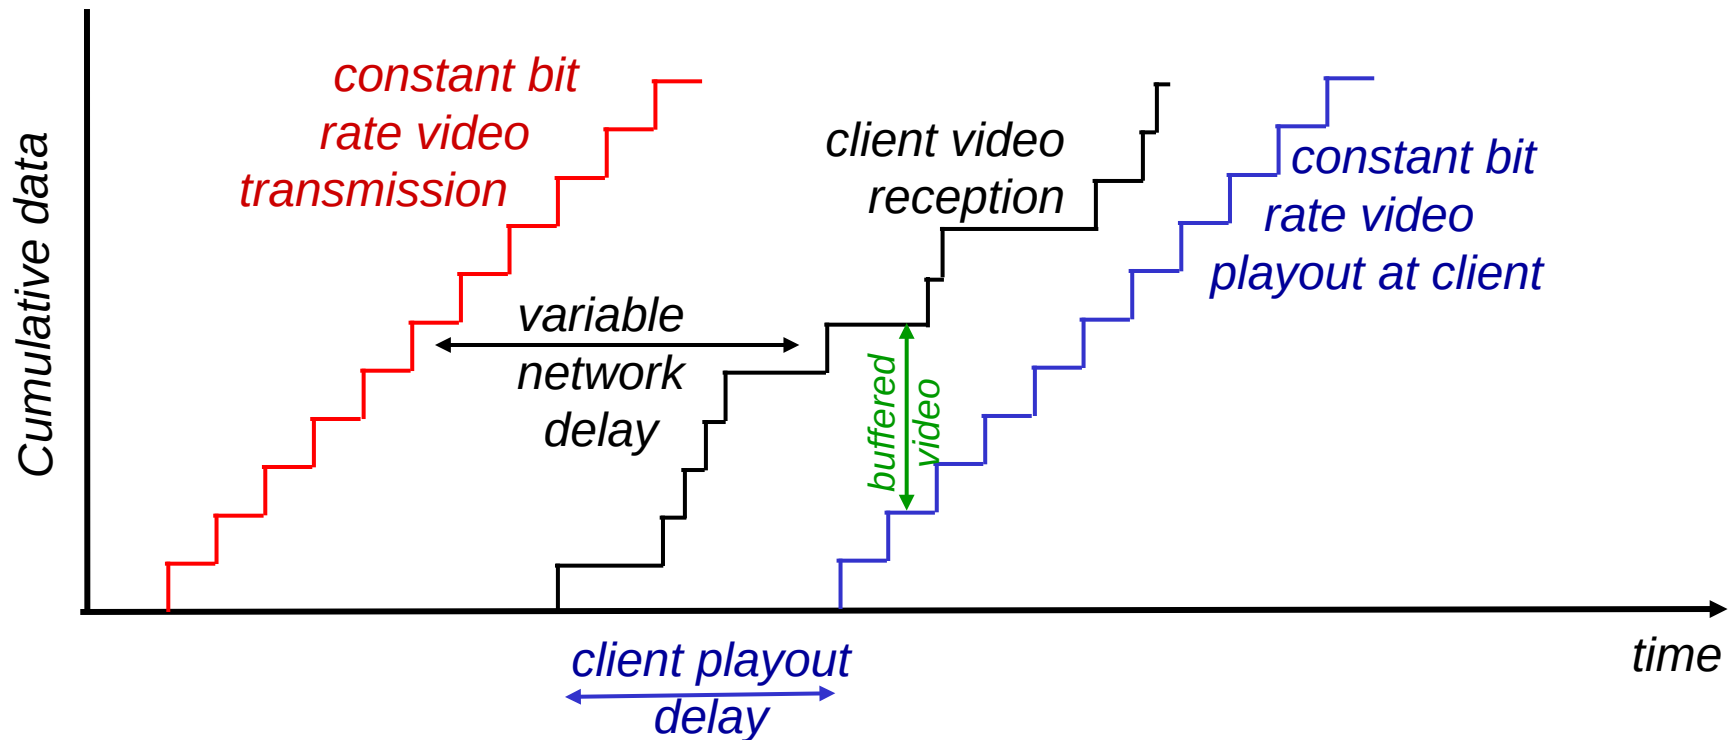
# Streaming stored video:



# Streaming stored video: challenges

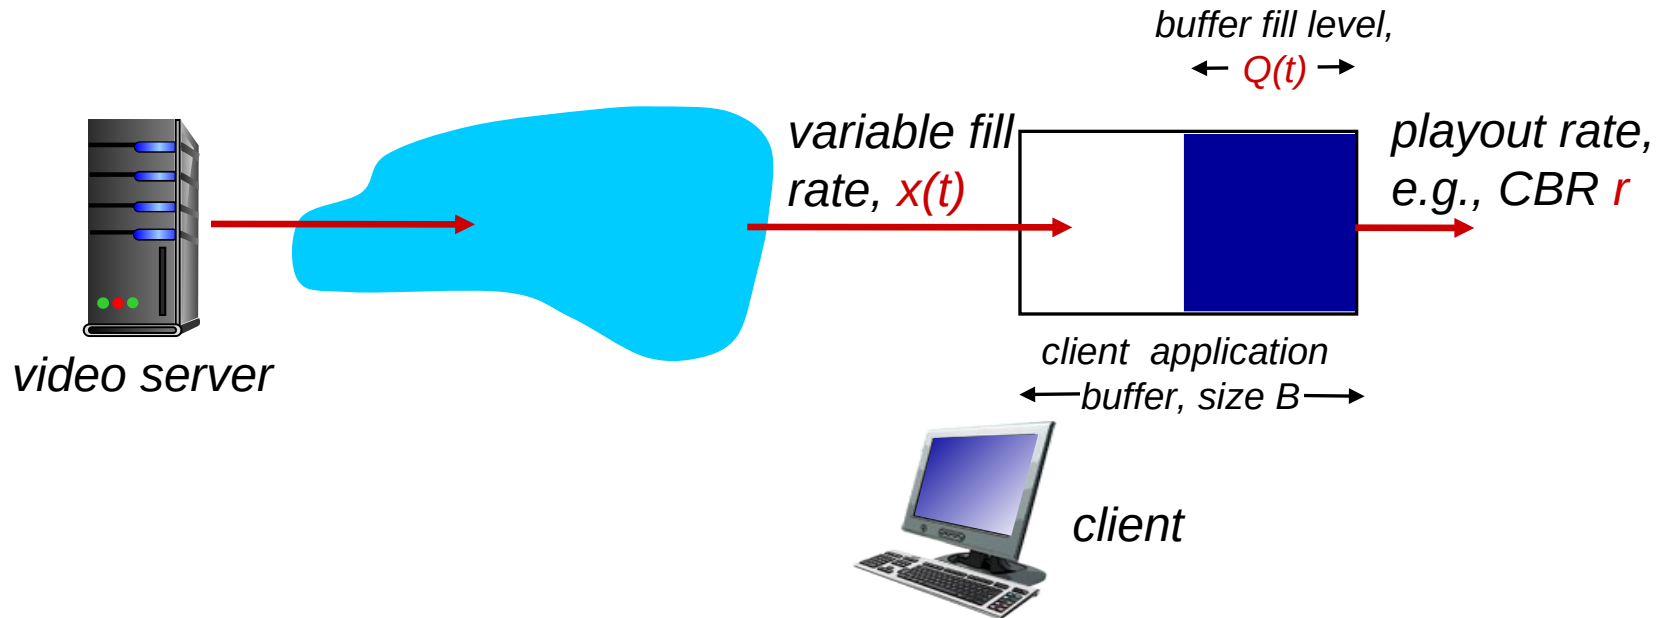
- ❖ *continuous playout constraint*: once client playout begins, playback must match original timing
  - ... but *network delays are variable* (jitter), so will need *client-side buffer* to match playout requirements
- ❖ *other challenges*:
  - *client interactivity*: pause, fast-forward, rewind, jump through video
  - *video packets may be lost, retransmitted*

# Streaming stored video: revisited



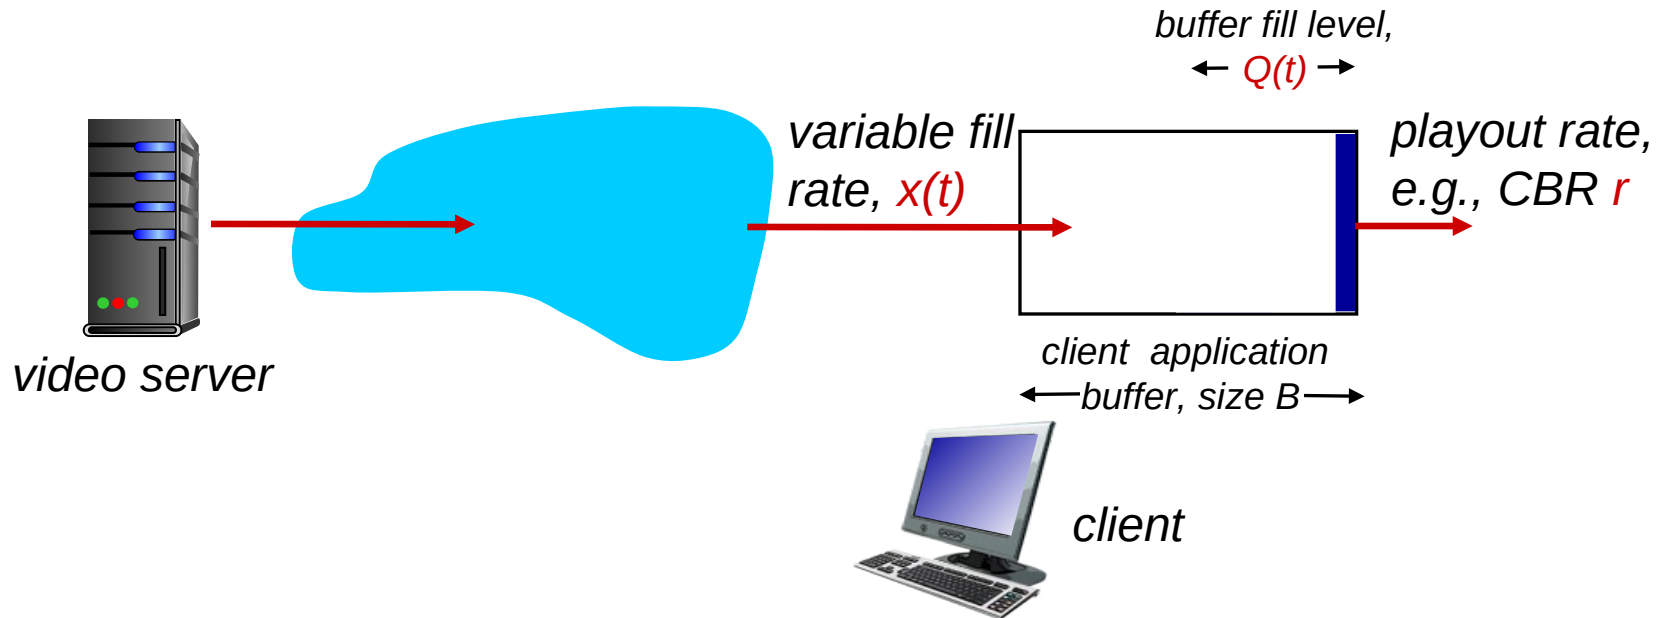
- ❖ *client-side buffering and playout delay:* compensate for network-added delay, delay jitter

# Client-side buffering, playout






# Client-side buffering, playout



1. Initial fill of buffer until playout begins at  $t_p$
2. playout begins at  $t_p$ ,
3. buffer fill level varies over time as fill rate  $x(t)$  varies and playout rate  $r$  is constant

# Client-side buffering, playout



*playout buffering: average fill rate ( $\bar{x}$ ), playout rate ( $r$ ):*

❖  $\bar{x} < r$ : buffer eventually empties (causing freezing of video playout until buffer again fills)

❖  $\bar{x} > r$ : buffer will not empty, provided initial playout delay is large enough to absorb variability in  $x(t)$


- *initial playout delay tradeoff*: buffer starvation less likely with larger delay, but larger delay until user begins watching

# Streaming multimedia: UDP

- ❖ *server sends at rate appropriate for client*
  - *often: send rate = encoding rate = constant rate*
  - *transmission rate can be oblivious to congestion levels*
- ❖ *short playout delay (2-5 seconds) to remove network jitter*
- ❖ *error recovery: application-level, time permitting*
- ❖ *RTP [RFC 2326]: multimedia payload types*
- ❖ *UDP may not go through firewalls*

# Streaming multimedia: HTTP

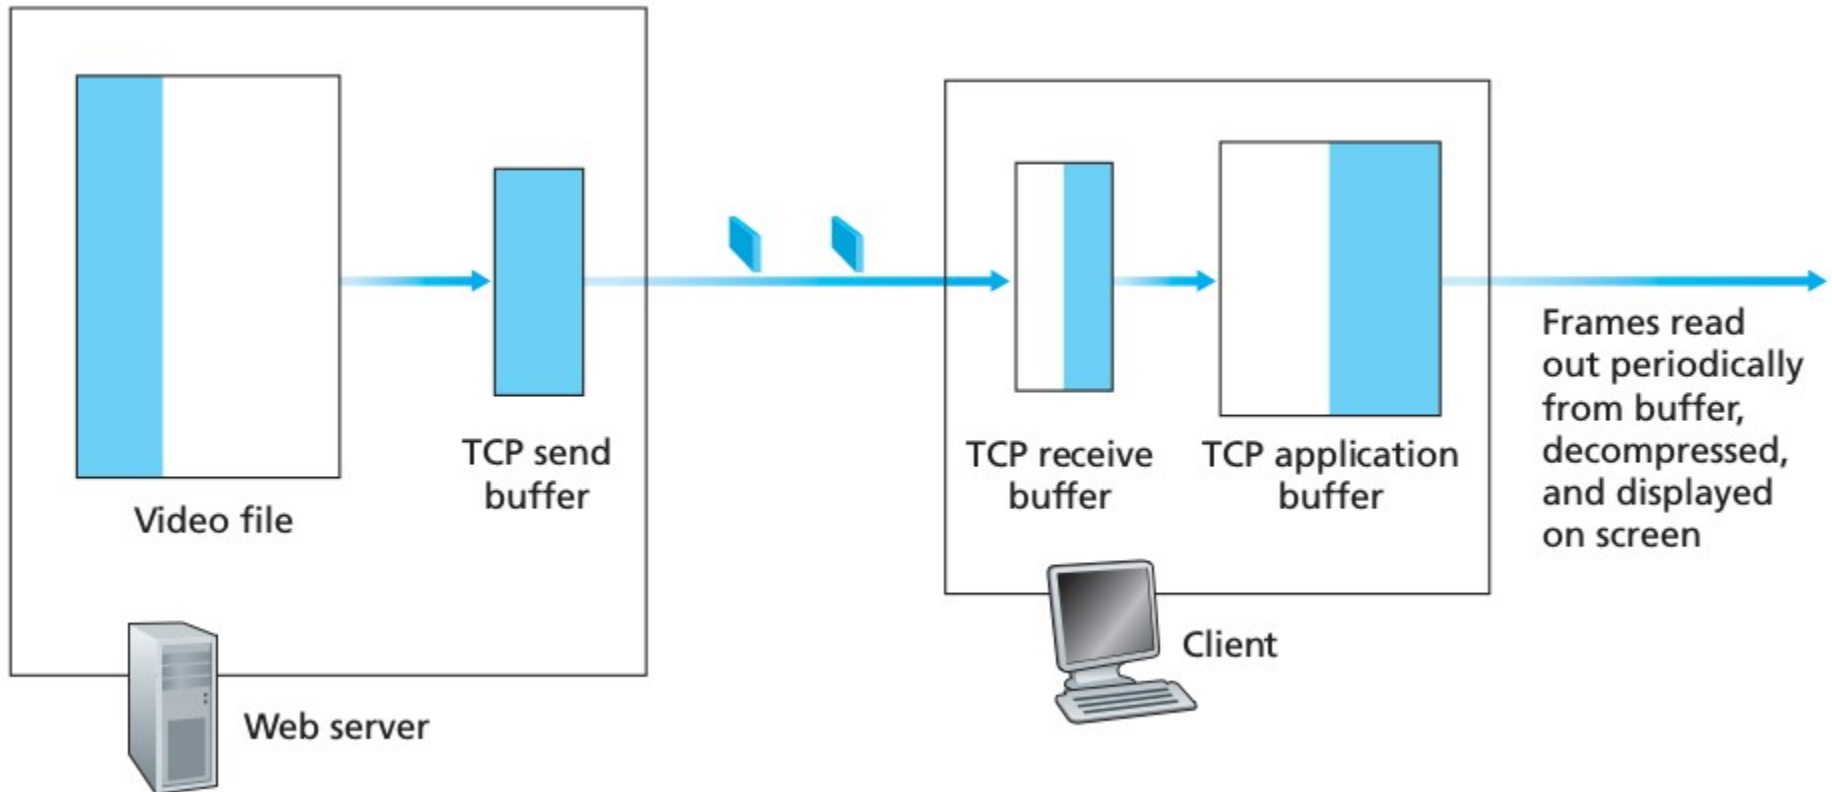
- ❖ multimedia file retrieved via HTTP GET
- ❖ send at maximum possible rate under TCP (if  $>$  consumption rate we have *prefetching*)



- ❖ fill rate fluctuates due to TCP congestion control, retransmissions (in-order delivery)
- ❖ larger playout delay: smooth TCP delivery rate
- ❖ HTTP/TCP passes more easily through firewalls

# Client Application and TCP Buffers

- ❖ *full client application buffer indirectly imposes limit on the transmission rate from server to client streaming over*



# Repositioning the Video

- ❖ *What happen when the user wants to jump to a future point in time in the video?*
- ❖ *Exploitation of the **HTTP byte-range header***
  - ❖ *specifies the specific range of bytes to retrieve*
- ❖ *User jump to a new position*
  - ❖ *client sends a new HTTP request with the byte-range header from which byte in the file should the server send data*
  - ❖ *the server receiving the new HTTP request can forget about any earlier request*
- ❖ *Jump to future point or earlier termination...*
  - ...waste of network bandwidth and server resources !*

# Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ *server:*
  - *divides video file into multiple chunks*
  - *each chunk stored, encoded at different rates*
  - *manifest file: provides URLs for different chunks*
- ❖ *client:*
  - *periodically measures server-to-client bandwidth*
  - *consulting manifest, requests one chunk at a time*
    - *chooses maximum coding rate sustainable given current bandwidth*
    - *can choose different coding rates at different points in time (depending on available bandwidth at time)*

# Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ “intelligence” at client: client determines
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



# Content distribution networks

*challenge: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?*

- ❖ *option 1: single, large “mega-server”*
  - *single point of failure*
  - *point of network congestion*
  - *long path to distant clients*
  - *multiple copies of video sent over outgoing link*

*....quite simply: this solution doesn't scale*

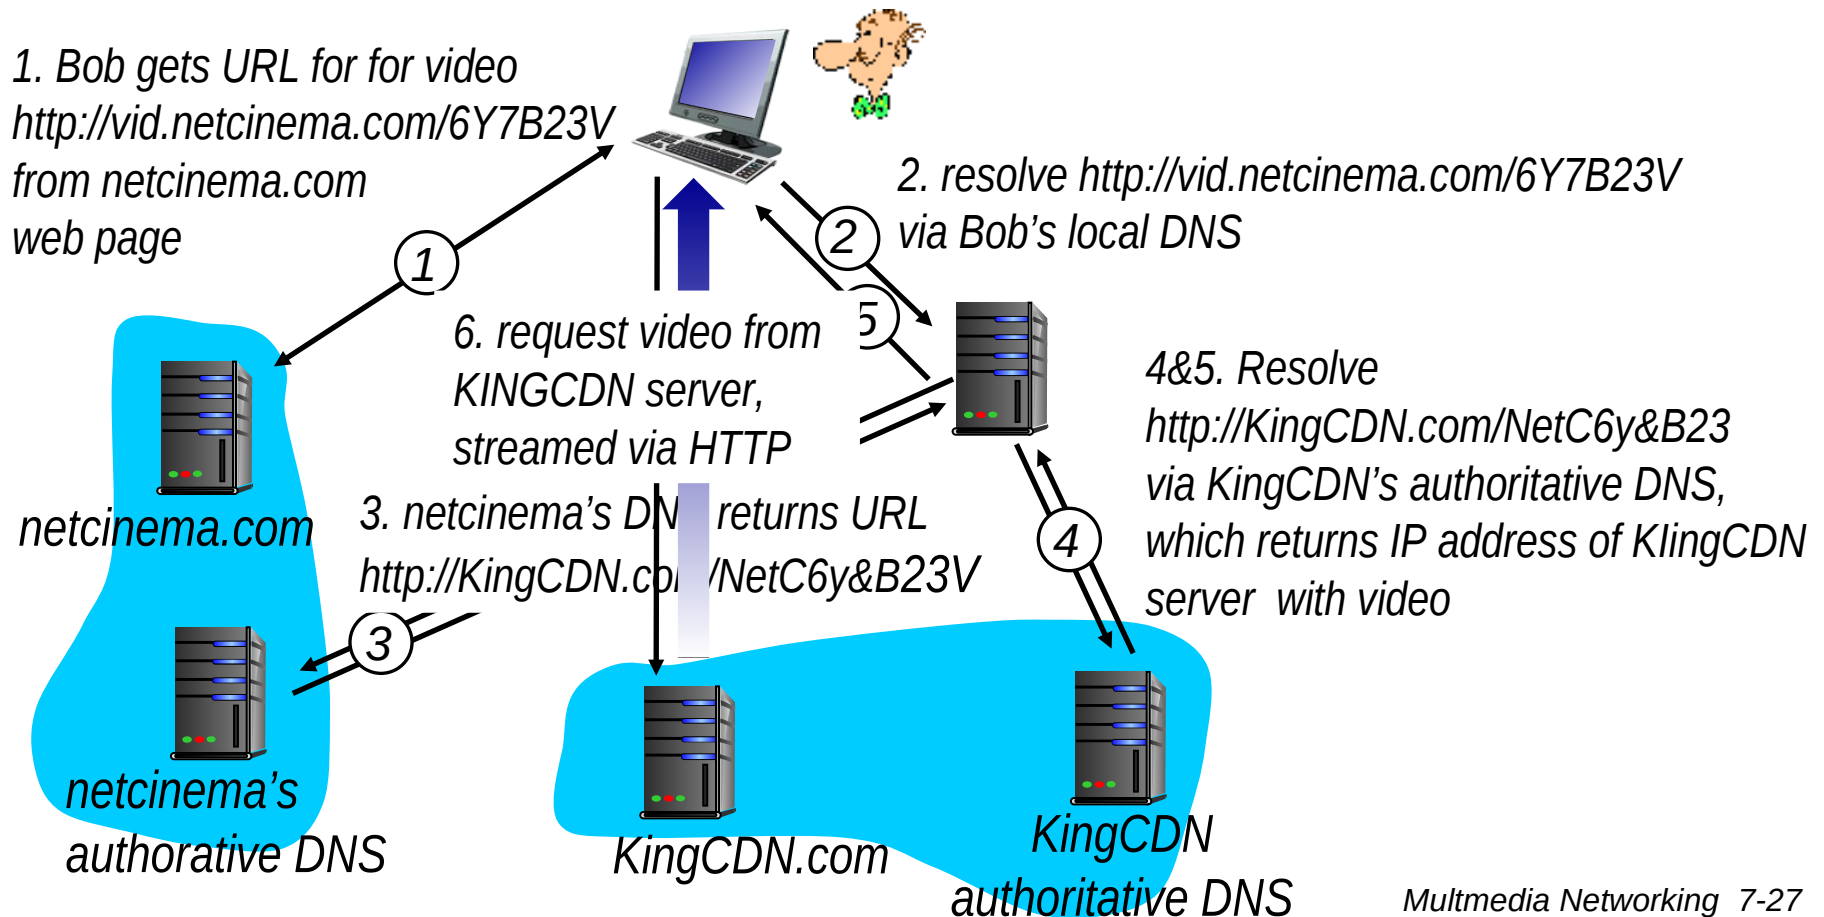
# Content distribution networks

- ❖ *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ❖ *option 2*: store/serve multiple copies of videos at multiple geographically distributed sites (CDN)
  - *enter deep*: push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - *bring home*: smaller number (10's) of larger clusters in key points near (but not within) access networks
    - used by Limelight

# CDN: “simple” content access scenario

Bob (client) requests video `http://vid.netcinema.com/6Y7B23V`

■ video stored in CDN at `http://KingCDN.com/NetC6y&B23V`



# CDN cluster selection strategy

- ❖ *challenge*: how does CDN DNS select “good” CDN node to stream to client
  - pick CDN node geographically closest to client
  - pick CDN node with shortest delay (or min # hops) to client (CDN nodes periodically ping access ISPs, reporting results to CDN DNS)
  - IP anycast
- ❖ *alternative*: let *client* decide - give client a list of several CDN servers
  - client pings servers, picks “best”
  - Netflix approach

# *Multimedia networking: outline*

*7.1 multimedia networking  
applications*

*7.2 streaming stored video*

*7.3 voice-over-IP*

# Voice-over-IP (VoIP)

- ❖ *VoIP end-end-delay requirement: needed to maintain “conversational” aspect*
  - *higher delays noticeable, impair interactivity*
  - *< 150 msec: good*
  - *> 400 msec bad*
  - *includes application-level (packetization, playout), network delays*
- ❖ *session initialization: how does callee advertise IP address, port number, encoding algorithms?*
- ❖ *value-added services: call forwarding, screening, recording*

# VoIP characteristics

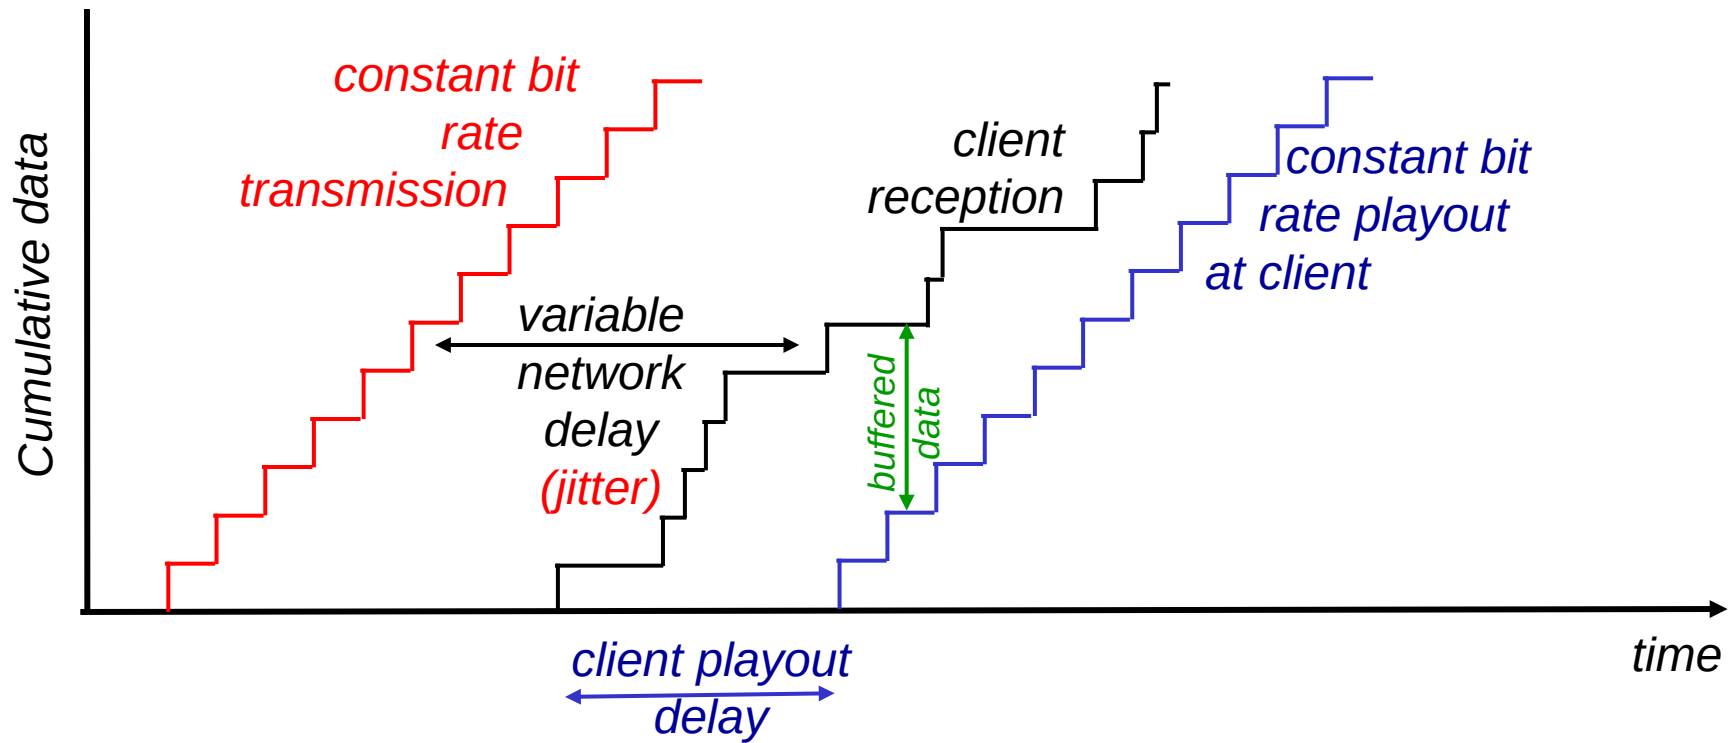
- ❖ *speaker's audio: alternating talk spurts, silent periods.*
  - *pkts generated only during talk spurts*
  - *Sender generates at 8 Kbytes/sec*
  - *With 20 msec chunks we have 160 bytes of data*
- ❖ *application-layer header added to each chunk*
- ❖ *chunk+header encapsulated into UDP or TCP segment*
- ❖ *application sends segment into socket every 20 msec during talkspurt*

# VoIP: packet loss, delay

- ❖ **network loss:** IP datagram lost due to network congestion (router buffer overflow)
- ❖ **delay loss:** IP datagram arrives too late for playout at receiver
  - delays: processing, queueing in network; end-system (sender, receiver) delays
  - typical maximum tolerable delay: 400 ms
- ❖ **loss tolerance:** depending on voice encoding, loss concealment, packet loss rates between 1% and 10% can be tolerated



# Delay jitter




- ❖ *end-to-end delays of two consecutive packets: difference can be more or less than 20 msec (transmission time difference)*

# VoIP: fixed playout delay

- ❖ *receiver attempts to playout each chunk exactly  $q$  msecs after chunk was generated.*
  - *chunk has time stamp  $t$ : play out chunk at  $t+q$*
  - *chunk arrives after  $t+q$ : data arrives too late for playout: data “lost”*
- ❖ *tradeoff in choosing  $q$ :*
  - *large  $q$ : less packet loss*
  - *small  $q$ : better interactive experience*

# VoIP: fixed playout delay

- sender generates packets every 20 msec during talk spurt.
- first packet received at time  $r$
- first playout schedule: begins at  $p$
- second playout schedule: begins at  $p'$



# Adaptive playout delay (1)

- ❖ *goal*: low playout delay, low late loss rate
- ❖ *approach*: adaptive playout delay adjustment:
  - estimate network delay, adjust playout delay at beginning of each talk spurt
  - silent periods compressed and elongated
  - chunks still played out every 20 msec during talk spurt
- ❖ adaptively estimate packet delay: (EWMA - exponentially weighted moving average, *recall TCP RTT estimate*):

$$d_i = (1-\alpha)d_{i-1} + \alpha (r_i - t_i)$$

delay estimate  
after  $i$ th packet

small constant,  
e.g. 0.1

time received - time sent  
(timestamp)  
measured delay of  $i$ th packet

## Adaptive playout delay (2)

- ❖ *also useful to estimate average deviation of delay,  $v_i$ :*

$$v_i = (1-\beta)v_{i-1} + \beta |r_i - t_i - d_i|$$

- ❖ *estimates  $d_i$ ,  $v_i$  calculated for every received packet, but used only at start of talk spurt*
- ❖ *for first packet in talk spurt, playout time is:*

$$\text{playout-time}_i = t_i + d_i + Kv_i$$

*remaining packets in talkspurt are played out periodically*

## Adaptive playout delay (3)

Q: How does receiver determine whether packet is first in a talkspurt?

- ❖ if no loss, receiver looks at successive timestamps
  - difference of successive stamps  $> 20$  msec --> talk spurt begins.
- ❖ with loss possible, receiver must look at both time stamps and sequence numbers
  - difference of successive stamps  $> 20$  msec *and* sequence numbers without gaps --> talk spurt begins.

# VoiP: recovery from packet loss (1)

**Challenge:** recover from packet loss given small tolerable delay between original transmission and playout

- ❖ each ACK/NAK takes  $\sim$  one RTT
- ❖ alternative: **Forward Error Correction (FEC)**
  - send enough bits to allow recovery without retransmission (recall two-dimensional parity in Ch. 5)


## **simple FEC**

- ❖ for every group of  $n$  chunks, create redundant chunk by exclusive OR-ing  $n$  original chunks
- ❖ send  $n+1$  chunks, increasing bandwidth by factor  $1/n$
- ❖ can reconstruct original  $n$  chunks if at most one lost chunk from  $n+1$  chunks, with playout delay

# VoiP: recovery from packet loss (2)

## *another FEC scheme:*

- ❖ *“piggyback lower quality stream”*
- ❖ *send lower resolution audio stream as redundant information*
- ❖ *e.g., nominal stream PCM at 64 kbps and redundant stream GSM at 13 kbps*



- ❖ *non-consecutive loss: receiver can conceal loss*
- ❖ *generalization: can also append (n-1)st and (n-2)nd low-bit rate chunk*



# VoiP: recovery from packet loss (3)




## *interleaving to conceal loss:*

- ❖ audio chunks divided into smaller units, e.g. four 5 msec units per 20 msec audio chunk
- ❖ packet contains small units from different chunks
- ❖ if packet lost, still have *most* of every original chunk
- ❖ no redundancy overhead, but increases playout delay

# Voice-over-IP: Skype


- ❖ *proprietary application-layer protocol (inferred via reverse engineering)*
  - *encrypted msgs*
- ❖ *P2P components:*
  - *clients:* skype peers connect directly to each other for VoIP call
  - *super nodes (SN):* skype peers with special functions
  - *overlay network:* among SNs to locate SCs
  - *login server*



# P2P voice-over-IP: skype

## *skype client operation:*

1. joins skype network by contacting SN (IP address cached) using TCP
2. logs-in (username, password) to centralized skype login server
3. obtains IP address for callee from SN, SN overlay
  - or client buddy list
4. initiate call directly to callee



# Skype: peers as relays

- ❖ **problem:** both Alice, Bob are behind “NATs”
  - NAT prevents outside peer from initiating connection to insider peer
  - inside peer can initiate connection to outside
- ❖ **relay solution:** Alice, Bob maintain open connection to their SNs
  - Alice signals her SN to connect to Bob
  - Alice’s SN connects to Bob’s SN
  - Bob’s SN connects to Bob over open connection Bob initially initiated to his SN

